

Zhaoyi Zhou

+86 18801383147 • zhouzhaoyi1123@gmail.com • <https://zhaoyizhou1123.github.io/>

EDUCATION

Institute for Interdisciplinary Information Sciences (IIIS), Tsinghua University

Expected June 2024

B.E. in Computer Science and Technology (Yao Class)

Beijing, China

- **Selected Coursework:** Machine Learning, Theory of Computation, Game Theory, Algorithm Design
- **GPA:** 3.97/4.00

PUBLICATIONS AND MANUSCRIPTS

- Convergence Rates for Localized Actor-Critic in Networked Markov Potential Games.
Zhaoyi Zhou, Zaiwei Chen, Yiheng Lin, and Adam Wierman.
Conference on Uncertainty in Artificial Intelligence (UAI) 2023. (31% acceptance rate)
- Free from Bellman Completeness: Trajectory Stitching via Model-based Return-conditioned Supervised Learning.
Zhaoyi Zhou, Chuning Zhu, Runlong Zhou, Qiwen Cui, Abhishek Gupta, and Simon S. Du.
Submitted to International Conference on Learning Representations (ICLR) 2024.

SELECTED HONORS

- Yao Award, Recognition Prize (presented by Dean Andrew C. Yao of IIIS), Tsinghua University, 2023;
- Scholarship of Academic Excellence, Tsinghua University, 2022;
- Scholarship of Scientific Innovation, Tsinghua University, 2022;
- Scholarship of Comprehensive Excellence, Tsinghua University, 2021.

RESEARCH EXPERIENCE

University of Washington

Feb. 2023 - Aug. 2023

Research Intern. Supervisor: Simon Shaolei Du

Seattle, USA

Led research in offline reinforcement learning (RL).

- Discovered freedom from Bellman completeness requirement as an advantage of return-conditioned supervised learning (RCSL) over dynamic-programming (DP) based offline RL methods.
- Conducted both theoretical and empirical analysis on explicit examples to demonstrate strength of RCSL over DP-based methods in near-deterministic environments given expert dataset.
- Theoretically prove that RCSL cannot do trajectory stitching, a limitation of RCSL observed in previous literature.
- Developed model-based return-conditioned supervised learning (MBRCSL) framework, which enables trajectory stitching of RCSL while avoiding Bellman completeness requirements.
- Tested MBRCSL on offline RL benchmarks, in which MBRCSL outperforms state-of-the-art model-based and model-free offline RL algorithms.

California Institute of Technology

Feb. 2022 - Feb. 2023

Research Intern. Supervisor: Adam Wierman

Remote

Led research in networked multi-agent RL.

- Introduced the class of networked Markov potential games as relaxation of Markov potential games (MPG), encompassing practical examples such as Markov congestion game.
- Designed localized actor-critic algorithm with linear function approximation, overcoming the curse of dimensionality.
- Derived the first finite-sample bound for multi-agent competitive games that is independent of the number of agents.
- Developed localized stochastic approximation (SA) framework as foundation of localized multi-agent policy evaluation.
- Derived finite-time convergence result of localized SA via a reduction approach.

PROJECT EXPERIENCE

Parameter-free Stochastic Optimization

June. 2023 - Present

Student Researcher, IIIS, Tsinghua University

Beijing, China

- Designed stochastic DoWG algorithm, a stochastic optimization method without manual parameter-tuning.
- Derived universal convergence rates of stochastic DoWG for smooth and non-smooth convex functions, matching optimal rate of tuned stochastic optimization methods up to an extra logarithm factor.
- Derive optimal step size for stochastic gradient descent (SGD) variants given strongly-convex objective functions.

OTHERS

Software: Python, C/C++, Go, Verilog, MATLAB

Languages: English (Fluent), German (Fluent), Chinese (Native)